

PROOFROUTE

ACCURACY FIRST. TOKENS LAST.



Not every query should pay the full price

ProofRoute turns inference cost into a decision.

Prove deterministic work locally. Use calibrated local intelligence only where evidence supports it. Buy remote intelligence when uncertainty demands it.

Accuracy is the gate. Token efficiency is the advantage.

ALWAYS REMOTE

FULL TOKEN COST

PROOFROUTE

LOCAL


REMOTE ONLY WHEN NEEDED

Two evidence-backed paths share one answer contract

SUBMITTED RUNTIME



CALIBRATED LOCAL CHALLENGER



The engine—not the model—owns JSON structure. Raw-prompt ablation cut Fireworks input tokens by 51.9% with byte-identical Kimi answers.

Deterministic answers require proof, not confidence

01 Parse a registered contract

bounded math, logic, extraction or code

02 Recompute the candidate

safe AST, graph closure or constrained execution

03 Require one unique result

ambiguity, mismatch or unsupported input refuses

04 Apply the answer contract

release zero Fireworks tokens or fall through

2 / 2

known correct ordering candidates recovered

5 / 5

incorrect deterministic candidates refused

Independent gate: 180 releases • 100%
precision • 0 false positives

AMD training turns 4,400 tasks into calibrated evidence

FunctionGemma scores the task. Gemma E2B answers. Mechanical proofs and independent judges label the outcome.



AMD ROCm 7.2 • PyTorch 2.9 • Q8 FunctionGemma • text-only LiteRT E2B

4,400 outputs; 2,621 mechanical verdicts; 1,779 semantic judgments.

Measured risk keeps local routing honest

2,183

correct E2B answers • 4,400-task corpus

95.95%

protected precision • best neural challenger

8.41%

protected coverage • 74 of 880

THE PUBLIC RELEASE

2.67 GB

compressed linux/amd64 image

712 TESTS

CI + Docker gates green

2 SEC

exact image resource gate

84.2% official control. Accuracy protected.